

Jev on MoCa: Agreement with Human Causal and Moral Judgments

September 30, 2026

Abstract

Across 206 English MoCa stories, Jev’s three-class agreement with the aggregated judgments of 25 human respondents per item is 43.69%. Agreement is 39.58% for causal items and 53.23% for moral items. Jev distinguishes clear human yes/no judgments but rarely places human-ambiguous causal stories in the ambiguous interval.

1 Dataset

The dataset includes 144 causal stories and 62 moral-permissibility stories, each with 25 binary human judgments. The classes are Yes, No and Ambiguous, using the 0.4/0.6 boundaries; Jev’s third class is derived from its binary-choice probability. Model snapshot: typesafe/jev-1.13-20260917.

Factor names follow the source’s expert annotations, and factors may overlap within a story. Human response proportions and Jev’s structured-choice probabilities have different meanings; this report describes their agreement and numerical differences on these stories.

2 Results

2.1 Overall results

Agreement is judged against the dataset’s aggregated human labels, without treating majority judgments as moral truth. Uniform three-class chance is 33.33%; the overall majority-class baseline is 36.41%. AUROC uses only human-unambiguous items.

Subset	Items	Agreement	95% interval	AUROC	Probability MAE
All	206	43.69%	36.89%–50.49%	0.704	0.308
Causal	144	39.58%	31.94%–47.92%	0.748	0.333
Moral	62	53.23%	40.32%–66.13%	0.852	0.249

2.2 Judgment distributions

Subset	Source	Yes	No	Ambiguous
Causal	Human	48	50	46
Causal	Jev	98	35	11
Moral	Human	23	10	29
Moral	Jev	22	21	19

Only 1 of 46 human-ambiguous causal items is also classified as ambiguous by Jev; the corresponding count is 12 of 29 moral items. Jev gives substantially more affirmative causal judgments, showing that lower three-class agreement involves distribution and boundary differences.

2.3 Mean differences across factor groups

Group transition	Group sizes	Human difference	Jev difference
Action to omission	80 / 32	+0.028	-0.179
Aware to unaware	23 / 20	-0.009	+0.195
Conjunctive to disjunctive	66 / 32	-0.030	-0.250
Abnormal to normal	44 / 45	-0.228	-0.294
Other to self beneficiary	26 / 22	+0.101	-0.015
Accidental to instrumental	18 / 30	-0.118	-0.067
Avoidable to inevitable	22 / 26	+0.131	+0.136
Impersonal to personal force	24 / 24	-0.058	-0.067

Differences describe changes in mean affirmative share/probability from the first group to the second. Factors overlap and story content changes across groups; these are descriptive associations, not isolated causal effects.

2.4 Call costs

Input tokens	Output tokens	Reported cost (USD)
109,940	6,798	0.00461748

All 206 requests succeeded on the first attempt, with no unknown charges or retries.

3 Conclusion

Jev shows partial agreement with aggregated human judgments, with higher agreement on moral than causal items. Its ability to distinguish clear human yes/no judgments does not translate into comparable identification of human ambiguity. Some factor-group mean directions match and others differ; the findings describe judgments on the same stories.

References

1. Nie et al. (2023). MoCa: Measuring Human-Language Model Alignment on Causal and Moral Judgment Tasks. NeurIPS. <https://arxiv.org/abs/2310.19677>
2. MoCa official repository, commit 1b61a20294247480d64675ceb19751ef4e1e878f. <https://github.com/cicl-stanford/moca>