

Jev on MoralChoice: Moral Choices and Response Stability

September 30, 2026

Abstract

Each of 1,367 English MoralChoice scenarios receives six structured judgments. All six variants select the reference action in every one of the 687 low-ambiguity scenarios. In the 680 high-ambiguity scenarios, 69.07% of judgments select action1 and 90.59% of scenarios retain the same underlying action across all variants. High-ambiguity choices are usually stable, with remaining sensitivity to form and order.

1 Dataset

The release contains 687 low-ambiguity and 680 high-ambiguity scenarios. The reference action is action1 for low-ambiguity scenarios; no universal correct action is assigned to high-ambiguity scenarios. Each scenario has ab, repeat and compare forms in both orders, yielding 8,202 judgments. Model snapshot: typesafe/jev-1.13-20260917.

Forms and order swaps follow the original design, while output is adapted to Jev’s structured binary choice. Six questions are answered in one request per scenario; repeat does not test verbatim text repetition. Order and cross-form stability are analysis measures for this adaptation, and six judgments are not counted as six independent scenarios.

2 Results

2.1 Overall results

Low-ambiguity judgments use the released reference action, with a 50% uniform binary-choice baseline. High-ambiguity results describe choices rather than accuracy.

Measure	Low ambiguity	High ambiguity
Scenarios	687	680
Judgments	4122	4080
Share selecting action1	100.00%	69.07%
Same action across all six	100.00%	90.59%
Mean P(action1)	99.82%	67.11%

Every low-ambiguity scenario selects the reference action in all variants: 100%, with a scenario-level Wilson 95% interval of 99.44%–100%. The high-ambiguity action1 share depends on the content arrangement in the release; it is not a first-position preference.

2.2 High-ambiguity stability

Form	Order agreement	Rate	Mean probability change	First-position share
ab	661 / 680	97.21%	3.15%	49.04%
repeat	658 / 680	96.76%	3.25%	48.97%
compare	643 / 680	94.56%	6.36%	47.57%

All six variants agree in 616 of 680 scenarios. The compare form has the lowest order agreement, indicating greater sensitivity of yes/no comparisons to order. The mean within-scenario probability range across six variants is 10.75 percentage points.

2.3 Auxiliary rule groups

Auxiliary rule labels come from the release and may overlap. Comparisons use scenarios where one action is marked as violating a rule and the other explicitly is not. Full aggregate tables are retained in the machine-readable summary.

Rule label	High-ambiguity scenarios	Nonviolating choice share	All six nonviolating
break_law	195	84.02%	76.92%
break_promise	108	68.21%	65.74%
cheat	138	93.24%	91.30%
death	84	64.88%	59.52%
deceive	262	86.39%	82.44%
disable	144	67.13%	61.11%
duty	409	86.15%	81.66%
freedom	219	54.26%	48.86%
pain	356	57.12%	51.97%
pleasure	276	50.60%	45.65%

2.4 Call costs

Input tokens	Output tokens	Reported cost (USD)
1,098,374	250,161	0.046131708

All 1,367 requests succeeded, with no retries or unknown charges.

3 Conclusion

Jev fully agrees with the reference choices in low-ambiguity scenarios. High-ambiguity scenarios have no universal correct answer; about nine in ten retain the same underlying action across the six form/order combinations, with more changes in the comparison form. The result describes structured choices and their stability, without equating concentrated probabilities with moral correctness or human opinion shares.

References

1. Scherrer et al. (2023). Evaluating the Moral Beliefs Encoded in LLMs. NeurIPS. https://proceedings.neurips.cc/paper_files/paper/2023/file/a2cf225ba392627529efef14dc857e22-Paper-Conference.pdf
2. MoralChoice data release, commit 89c0fe7b158b5ade5d10e0644c1aa20ab4c78cbe; CC BY 4.0. <https://huggingface.co/datasets/ninoscherrer/moralchoice>