

# Jev on SocialIQa: Social Commonsense Decisions

September 30, 2026

## Abstract

Jev answered 1,792 of 2,224 English SocialIQa v1.4 test questions correctly, achieving 80.58% accuracy with a duplicate-question cluster 95% interval of 78.81%–82.21%. Accuracy across the nine official source groups ranges from 78.11% to 84.66%. These results describe selection judgments about everyday motives, emotions and consequences.

## 1 Dataset

The material contains all 2,224 rows of the authors’ v1.4 test release. Each item contains an English context, a question and three options. Each is answered independently and scored against the released reference label. Model snapshot: typesafe/jev-1.13-20260917.

The nine groups use the official promptDim field, which records the ATOMIC question-generation source. Rewritten questions can depart from that source, so these groups are not treated as nine independent psychological abilities. The evaluation concerns a long-public English multiple-choice dataset.

## 2 Results

### 2.1 Overall results

The dataset labels are the scoring criterion. Every formal item has a valid response. Uniform random three-option selection has an expected accuracy of 33.33%.

| Measure                       | Result        |
|-------------------------------|---------------|
| Formal items                  | 2,224         |
| Correct                       | 1,792         |
| Accuracy                      | 80.58%        |
| Question-cluster 95% interval | 78.81%–82.21% |
| Distinct complete questions   | 2203          |
| Distinct-question accuracy    | 80.57%        |

## 2.2 Official source groups

| Group                       | Items | Correct | Accuracy |
|-----------------------------|-------|---------|----------|
| oEffect (Effect on others)  | 159   | 127     | 79.87%   |
| oReact (Others' reactions)  | 223   | 179     | 80.27%   |
| oWant (Others' wants)       | 252   | 201     | 79.76%   |
| xAttr (Actor attributes)    | 297   | 232     | 78.11%   |
| xEffect (Effect on actor)   | 103   | 82      | 79.61%   |
| xIntent (Actor intent)      | 297   | 239     | 80.47%   |
| xNeed (Actor prerequisites) | 277   | 225     | 81.23%   |
| xReact (Actor reactions)    | 277   | 220     | 79.42%   |
| xWant (Actor wants)         | 339   | 287     | 84.66%   |

The xWant group has the highest accuracy and xAttr the lowest. Group sizes differ; this ordering describes the dataset rather than a stable hierarchy of abilities.

## 2.3 Call costs

| Input tokens | Output tokens | Reported cost (USD) |
|--------------|---------------|---------------------|
| 851,327      | 84,512        | 0.035755734         |

Reported formal-evaluation cost is USD 0.035755734. Twelve separate interface-debugging items bring the reported total to USD 0.035949816, with 855,948 input tokens and 84,968 output tokens. One technical failure has unknown billing; the historical ledger retains USD 0.01 as a reservation, not a settled charge.

## 3 Conclusion

Jev achieves 80.58% accuracy on this English social-commonsense dataset, selecting the reference answer for most questions about motives, emotions and consequences. It misses 432 reference answers. The result provides a baseline for social-situation understanding under the tested selection format; it does not measure emotional experience or willingness to help.

## References

1. Sap et al. (2019). Social IQa: Commonsense Reasoning about Social Interactions. EMNLP-IJCNLP. <https://aclanthology.org/D19-1454/>
2. SocialIQA v1.4: authors' dataset release, CC BY 4.0. <https://maartensap.com/social-iqa/>